

Jiwon Cha  
jc4va

## In-lab report

### Topic

Optimized code: Compare code generated normally to optimized code. To create optimized code, you will need to use the -O2 compiler flag. Can you make any guesses as to why the optimized code looks as it does? What is being optimized? Be sure to show your original sample code as well as the optimized version. Try loops and function calls to see what "optimizing" does. Be aware that if instructions are "not necessary" to the final output of the program then they may be optimized away completely! This does not lead to very interesting comparisons. Describe at least four (non-trivial) differences you see between 'normal' code and optimized code.

### Code 1

C++ code:

- `int foo(int a){return a+3};`
- `int main(){foo(3)};`

Assembly code – without –O2 flag:

```
mov     dword ptr [rbp - 4], edi
mov     edi, dword ptr [rbp - 4]
add     edi, 3
mov     eax, edi
pop     rbp
ret

sub     rsp, 16
mov     edi, 3
call    __Z3fooi
xor     edi, edi
mov     dword ptr [rbp - 4], eax ## 4-byte Spill
mov     eax, edi
add     rsp, 16
pop     rbp
ret
.cfi_endproc
```

Assembly code – with –O2 flag:

```
lea     eax, [rdi + 3]
pop     rbp
ret

.cfi_def_cfa_register rbp
xor     eax, eax
pop     rbp
ret
```

Foo function:

First of all, it is readily apparent that the optimized version is much shorter than the original assembly code. The optimized version ignores several unnecessary steps such as moving the element in edi to a spot on the stack and then moving the element in that spot into edi again. This step is not necessary within the context of this code, so the assembly takes it out. It also uses lea instruction instead of mov, which makes sense because lea can add and load at the same time while mov requires separate instruction for load and add. For example, if you want to add several constants to the register, lea is much shorter and faster. The code above shows how several steps can be shortened to one instruction; setting up edi, add edi, 3; `mov eax, edi` → `lea eax, [rdi+3]`

Main method:

Main method is also very short compared to the version without –O2 flag. First of all, the original main method manually adds 3 to the first parameter and then moves that value to eax. It also performs additional operations such as moving rsp to the original spot. However, the optimized version does not need to perform this step because 3 was already added to the first parameter inside of foo function. Therefore, in the main method, it simply returns what is in rax by calling ret.

## Code 2

Assembly code for this program only includes the function part (foo) as the main function is the same as code 1.

### C++ code:

```
int foo(int a){
    for (int i = 0; i < 5; i++){
        a += i;
    }
    return a;
}

int main(){
    foo(3);
}
```

### Assembly code – without –O2 flag:

```
Ltmp2:
    .cfi_def_cfa_register rbp
    mov     dword ptr [rbp - 4], 0
    mov     dword ptr [rbp - 12], 0
LBB0_1:
    cmp     dword ptr [rbp - 12], 5      ## =>This Inner Loop Header: Depth=1
    jge     LBB0_4
## BB#2:
    mov     eax, dword ptr [rbp - 12]    ## in Loop: Header=BB0_1 Depth=1
    add     eax, dword ptr [rbp - 8]
    mov     dword ptr [rbp - 8], eax
## BB#3:
    mov     eax, dword ptr [rbp - 12]    ## in Loop: Header=BB0_1 Depth=1
    add     eax, 1
    mov     dword ptr [rbp - 12], eax
    jmp     LBB0_1
LBB0_4:
    mov     eax, dword ptr [rbp - 4]
    pop     rbp
    ret
    .cfi_endproc
```

### Assembly code – with –O2 flag:

```
.cfi_def_cfa_register rbp

lea     eax, [rdi + 10]
pop     rbp
ret
.cfi_endproc
```

The original version of assembly carries out each step of the loop. We see that arguments of the for loop (such as int i) are stored on the stack as local variables. int i starts from 0 and increments until it reaches 5, therefore it first moves 0 to a spot on the stack and then compares the value stored in that spot to 5. If it is equal to 5, it jumps to the last part of loop and returns the value in eax. Otherwise, it keeps adding the value of i to eax and then incrementing the value in spot [rbp-12].

However, the optimized version skips most of the steps described above. The function simply adds 10 to the parameter and moves that value to eax, which means it has already calculated the result of 0+1+2+3+4. This is very convenient because the method is simply adding 0, 1, 2, 3, and 4 to the parameter.

To conclude, the optimized version seems to skip many steps unlike the original versions, and sometimes take parameters passed in the main function directly in the callee (foo function); it carries out operations that are usually performed in the main method, inside of the callee. The optimized version also produces significantly shorter code (less instructions) and less time consuming operations.

## Topic

Dynamic dispatch: Describe how dynamic dispatch is implemented. Note that dynamic dispatch is NOT the same thing as dynamic memory! Show this using a simple class hierarchy that includes virtual functions. Use more than one virtual function per class.

### C++ code and dynamic dispatch

```
class A{
public:
    virtual int add(int x){return x+30;};
    virtual int subtract(int x){return x-70;};
    int multiply(int x){return x*1;};
    int remainder (int x){return x;};
};

class B: public A{
public:
    virtual int add(int x){return x+60;};
    int subtract(int x){return x-140;};
    virtual int multiply(int x){return x*2;};
    int remainder (int x){return x*2;};
};

int main(){
    A *hi = new B();
    B *hi2 = new B();
    hi->add(0);
    hi->subtract(0);
    hi->multiply(1) << endl;
    hi->remainder(2);
}
```

#### Output:

```
60 (B)
-140 (B)
1 (A)
2 (A)
```

**Dynamic dispatch:** the process of selecting which implementation of a polymorphic operation (method or function) to call at run time.<sup>1</sup>

In C++, dynamic dispatch is performed using the 'virtual' keyword. "virtual" allows programmers to obtain 'late binding' where which implementation of the method is used is decided at compile time based on the type of the pointer that he calls through.<sup>2</sup>

From the code and output above, we see that the program uses the function of superclass when the methods are declared without 'virtual'; hi->remainder(2) yields 2, which means it uses the remainder function of 'A' class. However, we see that the program uses add and subtract method of B classes when we put virtual keyword in front of method names in 'A' class.

---

<sup>1</sup> [https://en.wikipedia.org/wiki/Dynamic\\_dispatch](https://en.wikipedia.org/wiki/Dynamic_dispatch)

<sup>2</sup> <https://stackoverflow.com/questions/2391679/why-do-we-need-virtual-functions-in-c>

## Assembly code and implementation

### Add & subtract

```
.cfa_def_cfa_register rbp
sub    rsp, 32
mov    eax, 8
mov    edi, eax
call   __Znwm
xor    esi, esi
mov    ecx, 8
mov    edx, ecx
mov    rdi, rax
mov    qword ptr [rbp - 16], rax ## 8-byte Spill
call   _memset
mov    rdi, qword ptr [rbp - 16] ## 8-byte Reload
call   __ZN1BC1Ev
xor    esi, esi
mov    rax, qword ptr [rbp - 16] ## 8-byte Reload
mov    qword ptr [rbp - 8], rax
mov    rax, qword ptr [rbp - 8]
mov    rdx, qword ptr [rax]
mov    rdi, rax
call   qword ptr [rdx]
xor    ecx, ecx
mov    dword ptr [rbp - 20], eax ## 4-byte Spill
mov    eax, ecx
add    rsp, 32
pop    rbp
ret
.cfi_endproc
```

### multiply

```
.cfa_def_cfa_register rbp
sub    rsp, 32
mov    eax, 8
mov    edi, eax
call   __Znwm
xor    esi, esi
mov    ecx, 8
mov    edx, ecx
mov    rdi, rax
mov    qword ptr [rbp - 16], rax ## 8-byte Spill
call   _memset
mov    rdi, qword ptr [rbp - 16] ## 8-byte Reload
call   __ZN1BC1Ev
mov    esi, 1
mov    rax, qword ptr [rbp - 16] ## 8-byte Reload
mov    qword ptr [rbp - 8], rax
mov    rdi, qword ptr [rbp - 8]
call   __ZN1A8multiplyEi
xor    ecx, ecx
mov    dword ptr [rbp - 20], eax ## 4-byte Spill
mov    eax, ecx
add    rsp, 32
pop    rbp
ret
.cfi_endproc
```

\*Note that add and subtract use the member functions of class B (virtual dispatch) and multiply uses the member function of class A (static dispatch).

The general syntax for the methods for add(subtract) and multiply are very similar. Both function uses 'hidden parameter register' such as ecx and edx as it calls the functions.<sup>3</sup> They also both use `call _memset`, which sets the first number of bytes of the block of memory pointed by ptr to a specific spot.<sup>4</sup>

The main difference that we see is that 'multiply' function calls the function directly via the instruction `call __ZN1A8multiplyEi`, which indeed calls the multiply function of 'A'. However, virtual dispatch 'loads the address of the vtable from the object, then loads the function address from the vtable, then finally calls the function indirectly'.<sup>5</sup> Indeed, we see that the assembler loads the address at [rbp-8] to rax, then moves that address to rdx, then finally calls the pointer to [rdx]. This process facilitates dynamic binding.

Even though the syntax is rather similar, dynamic dispatch is different from dynamic memory. Dynamic dispatch involves selecting which member function or method to call (polymorphic implementation). They are similar in that dynamic memory allocation and dynamic dispatch both involve run-time operation.

<sup>3</sup> <http://loci-lang.org/DynamicDispatch.html>

<sup>4</sup> [www.cplusplus.com/reference/cstring/memset/](http://www.cplusplus.com/reference/cstring/memset/)

<sup>5</sup> <https://stackoverflow.com/questions/20147054/how-does-dynamic-dispatch-happen-in-assembly>

**References:**

[https://groups.google.com/forum/#!topic/comp.lang.asm.x86/HhOsbmwyU\\_I](https://groups.google.com/forum/#!topic/comp.lang.asm.x86/HhOsbmwyU_I)

<https://stackoverflow.com/questions/20147054/how-does-dynamic-dispatch-happen-in-assembly>

<https://stackoverflow.com/questions/9995922/how-to-tell-if-a-program-uses-dynamic-dispatch-by-looking-at-the-assembly>

<http://loci-lang.org/DynamicDispatch.html>

<https://stackoverflow.com/questions/2391679/why-do-we-need-virtual-functions-in-c>